

INF-2201

06 – Semaphores

John Markus Bjørndalen
2025

Synchronization recap

- No preemption -- relatively easy to get a decent solution:
 - Can turn off interrupts to turn a preemptive to a nonpreemptive environment
 - Move synchronization to the kernel if necessary
- If preemption:
 - A correct solution that doesn't use hardware support quickly gets a) complicated and b) hard to prove and c) use spinning
 - Solutions that use hardware features (see figures on the right) are easier to reason about, but can still waste cycles while spinning
 - Priority inversion can be an issue

enter_region:

```
TSL REGISTER,LOCK
CMP REGISTER,#0
JNE enter_region
RET
```

I copy lock to register and set lock to 1
I was lock zero?
I if it was not zero, lock was set, so loop
I return to caller; critical region entered

leave_region:

```
MOVE LOCK,#0
RET
```

I store a 0 in lock
I return to caller

Figure 2-25. Entering and leaving a critical section using the TSL instruction.
Modern Operating Systems, Tanenbaum & Bos

enter_region:

```
MOVE REGISTER,#1
XCHG REGISTER,LOCK
CMP REGISTER,#0
JNE enter_region
RET
```

I put a 1 in the register
I swap the contents of the register and lock variable
I was lock zero?
I if it was non zero, lock was set, so loop
I return to caller; critical region entered

leave_region:


```
MOVE LOCK,#0
RET
```

I store a 0 in lock
I return to caller

Figure 2-26. Entering and leaving a critical section using the XCHG instruction.
Modern Operating Systems, Tanenbaum & Bos

Easy solution 1 (for a single core OS)

- Provide a `enter_region` system call that
 - Enters the kernel
 - Disables interrupts
 - Runs the "try part" of `enter_region` on the right (don't use the loop)
 - If it doesn't succeed, blocks the thread/process and puts it on a waiting queue. Then run scheduler
 - If it succeeds, enable interrupts and return
- Provide a `leave_region` system call that
 - Enters the kernel
 - Disables the interrupts
 - Pops out one or more processes/threads from the waiting queue
 - Runs the `leave_region` bit on the right
 - Enables interrupts and returns
- High overhead



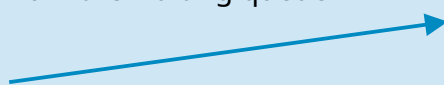
```
enter_region:
    TSL REGISTER,LOCK
    CMP REGISTER,#0
    JNE enter_region
    RET
```

I copy lock to register and set lock to 1
I was lock zero?
I if it was not zero, lock was set, so loop
I return to caller; critical region entered

```
leave_region:
    MOVE LOCK,#0
    RET
```

I store a 0 in lock
I return to caller

Figure 2-25. Entering and leaving a critical section using the TSL instruction.
Modern Operating Systems, Tanenbaum & Bos



```
enter_region:
    MOVE REGISTER,#1
    XCHG REGISTER,LOCK
    CMP REGISTER,#0
    JNE enter_region
    RET
```

I put a 1 in the register
I swap the contents of the register and lock variable
I was lock zero?
I if it was non zero, lock was set, so loop
I return to caller; critical region entered

```
leave_region:
    MOVE LOCK,#0
    RET
```

I store a 0 in lock
I return to caller

Figure 2-26. Entering and leaving a critical section using the XCHG instruction.
Modern Operating Systems, Tanenbaum & Bos

Producer-consumer problem


The producer-consumer problem has the following parts

- 1) A producer: creates information that it adds to a queue. If the queue is full, it waits until there is room.
- 2) A queue (can be a circular buffer) with a limited number of slots (bounded buffer)
- 3) A consumer: waits until there is something in the queue, pulls the first item out and consumes it (ignore how).

The solutions has to observe the following:

- The producer and consumer are executing independently. You cannot assume anything about their speeds.
- The queue data structure must be preserved. A simple solution is to use a critical region around anything that handles the queue.

More general case: there might be multiple producers and consumers, but we will ignore that for now.



Producer-consumer – solution 1

- Solution that assumes only one producer and one consumer
- Both use sleep if they cannot continue adding or removing items from the queue
- The other end uses wakeup if they add or remove items
- Race condition... where?

```
#define N 100                                /* number of slots in the buffer */
int count = 0;                               /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                            /* repeat forever */
        item = produce_item();               /* generate next item */
        if (count == N) sleep();             /* if buffer is full, go to sleep */
        insert_item(item);                  /* put item in buffer */
        count = count + 1;                   /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);    /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                            /* repeat forever */
        if (count == 0) sleep();             /* if buffer is empty, got to sleep */
        item = remove_item();               /* take item out of buffer */
        count = count - 1;                   /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer); /* was buffer full? */
        consume_item(item);                 /* print item */
    }
}
```

Figure 2-27. Producer-consumer problem with a fatal race condition.
Modern Operating Systems, Tanenbaum & Bos

Producer-consumer – solution 1

- Race condition... where?
- A general tool for spotting potential race conditions
 - These are equivalent for this discussion
 - Read-modify-write
 - Observe-decide-act
 - If the three are not done atomically (anybody can modify state between the steps), then there is a chance of a race condition

```
#define N 100                                /* number of slots in the buffer */
int count = 0;                               /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                            /* repeat forever */
        item = produce_item();                /* generate next item */
        if (count == N) sleep();               /* if buffer is full, go to sleep */
        insert_item(item);                    /* put item in buffer */
        count = count + 1;                     /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);     /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                            /* repeat forever */
        if (count == 0) sleep();               /* if buffer is empty, got to sleep */
        item = remove_item();                  /* take item out of buffer */
        count = count - 1;                     /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer); /* was buffer full? */
        consume_item(item);                    /* print item */
    }
}
```

Figure 2-27. Producer-consumer problem with a fatal race condition.
Modern Operating Systems, Tanenbaum & Bos

Producer-consumer – solution 1

- Using the tool 1 (producer vs. consumer):
 - "if (count == N) sleep()" is an **observe** (read count), **decide** (compare count to N), and **act** (sleep).
 - There is no protection of the state, so the consumer may remove an item between reading count and the decide step in the producer.
 - The producer is not woken up even if there is room. May wake up the next time the consumer removes an item.
- Using the tool 2 (consumer vs. producer):
 - Observe from the point of the consumer
 - If consumer preempted just before sleep, it might miss a wakeup from producer (count 0->1). Since this is the only chance, consumer never wakes up.

```
#define N 100                                /* number of slots in the buffer */
int count = 0;                               /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                            /* repeat forever */
        item = produce_item();               /* generate next item */
        if (count == N) sleep();             /* if buffer is full, go to sleep */
        insert_item(item);                  /* put item in buffer */
        count = count + 1;                   /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);    /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                            /* repeat forever */
        if (count == 0) sleep();             /* if buffer is empty, got to sleep */
        item = remove_item();               /* take item out of buffer */
        count = count - 1;                   /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer); /* was buffer full? */
        consume_item(item);                 /* print item */
    }
}
```

Figure 2-27. Producer-consumer problem with a fatal race condition.
Modern Operating Systems, Tanenbaum & Bos

Semaphores

- Locks typically have two states:
 - 0: lock free / released
 - 1: lock taken / acquired
- A more general concept is a semaphore
 - General idea: use an integer to store the number of wakeups. Can be larger than 1!
 - Two operations (similar to acquire and release). Both are atomic:
 - **Down:** check if value is larger than 0. If it was 0, sleep / wait. When it is 1 or larger: count down by one.
 - **Up:** add one to the semaphore. If there is a waiting process, release it.
- Can make user level / spinning semaphores
 - Same problem as with locks etc
- More useful if the waiting process is blocked and put on a queue (atomically) with the help of the operating system.

Pthreads semaphores

```
#include <semaphore.h>

int sem_init(sem_t *sem, int pshared,
             unsigned int value);

int sem_post(sem_t *sem);

int sem_wait(sem_t *sem);

int sem_trywait(sem_t *sem);

int sem_timedwait(sem_t *sem,
                  const struct timespec *abs_timeout);
```

From manpage of `sem_init`, `sem_post`, `sem_wait`

Semaphores

- Solving producer-consumer using semaphores
- To understand how it works:
 - The mutex semaphore is used to protect the queue datastructure
 - Note empty=N and full=0
 - Try looking at two cases first
 - 1) Producer running until the queue is full
 - 2) Consumer running until it blocks

```
#define N 100                                     /* number of slots in the buffer */
typedef int semaphore;                             /* semaphores are a special kind of int */
semaphore mutex = 1;                               /* controls access to critical region */
semaphore empty = N;                               /* counts empty buffer slots */
semaphore full = 0;                                /* counts full buffer slots */

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();                    /* TRUE is the constant 1 */
        down(&empty);                             /* generate something to put in buffer */
        down(&mutex);                             /* decrement empty count */
        insert_item(item);                         /* enter critical region */
        up(&mutex);                                /* put new item in buffer */
        up(&full);                                 /* leave critical region */
        up(&full);                                 /* increment count of full slots */
    }
}



void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);                               /* infinite loop */
        down(&mutex);                             /* decrement full count */
        item = remove_item();                      /* enter critical region */
        up(&mutex);                                /* take item from buffer */
        up(&empty);                                /* leave critical region */
        consume_item(item);                       /* increment count of empty slots */
        do_something_with_item(item);              /* do something with the item */
    }
}
```

Figure 2-28. The producer-consumer problem using semaphores.
Modern Operating Systems, Tanenbaum & Bos

Readers and writers problem



- $\#r \geq 1$ readers trying to read state
- $\#w \geq 1$ writers trying to write state
- Special cases where $\#r$ or $\#w$ is 1.
- Rules:
 - If no writers are active, then multiple readers can be active at the same time (no changes to the state)
 - If a writer is changing state, then we should not allow any readers (observe inconsistent state) or any other writers (inconsistent updates to state)
- Can get better performance as multiple readers can be serviced at the same time.
- Example from the book: airline reservation database.



Readers and writers problem

Need to keep track of the following invariants

- Number of readers ($\#r$):
 - $0..R$ if $\#w == 0$, 0 if $\#w > 0$
- Number of writers:
 - $0..1$ if $\#r == 0$, 0 if $\#r > 0$



One solution using semaphores

Note the asymmetry

- Writer(s) lock the entire db so only one writer can be inside at any point in time, and it also locks out readers
- The first reader locks down the database on behalf of other readers
 - Releases the mutex semaphore to let other readers in
 - The last reader to exit releases the db
- Note: there is a fairness issue with this solution. A continuous stream of readers may keep the writers locked out indefinitely.

```
typedef int semaphore;          /* use your imagination */
semaphore mutex = 1;           /* controls access to rc */
semaphore db = 1;              /* controls access to the database */
int rc = 0;                    /* # of processes reading or wanting to */

void reader(void)
{
    while (TRUE) {              /* repeat forever */
        down(&mutex);           /* get exclusive access to rc */
        rc = rc + 1;            /* one reader more now */
        if (rc == 1) down(&db); /* if this is the first reader ... */
        up(&mutex);             /* release exclusive access to rc */
        read_data_base();       /* access the data */
        down(&mutex);           /* get exclusive access to rc */
        rc = rc - 1;            /* one reader fewer now */
        if (rc == 0) up(&db);   /* if this is the last reader ... */
        up(&mutex);             /* release exclusive access to rc */
        use_data_read();        /* noncritical region */
    }
}

void writer(void)
{
    while (TRUE) {              /* repeat forever */
        think_up_data();        /* noncritical region */
        down(&db);              /* get exclusive access */
        write_data_base();      /* update the data */
        up(&db);                /* release exclusive access */
    }
}
```

Figure 2-29. A solution to the readers and writers problem.
Modern Operating Systems, Tanenbaum & Bos