

INF-2201

06 – Semaphores

John Markus Bjørndalen
2025

Synchronization recap

- No preemption -- relatively easy to get a decent solution:

- Can turn off interrupts to turn a preemptive to a nonpreemptive environment
- Move synchronization to the kernel if necessary

- If preemption:

- A correct solution that doesn't use hardware support quickly gets a) complicated and b) hard to prove and c) use spinning
- Solutions that use hardware features (see figures on the right) are easier to reason about, but can still waste cycles while spinning
- Priority inversion can be an issue

```
enter_region:
    TSL REGISTER, LOCK
    CMP REGISTER, #0
    JNE enter_region
    RET

leave_region:
    MOV LOCK, #0
    RET
```

I copy lock to register and set lock to 1
I was lock zero?
If it was not zero, lock was set, so loop
I return to caller; critical region entered

I store a 0 in lock
I return to caller

Figure 2-25. Entering and leaving a critical section using the TSL instruction.
Modern Operating Systems, Tanenbaum & Bos

```
enter_region:
    MOV REGISTER, #1
    XCHG REGISTER, LOCK
    CMP REGISTER, #0
    JNE enter_region
    RET

leave_region:
    MOV LOCK, #0
    RET
```

I put a 1 in the register
I swap the contents of the register and lock variable
I was lock zero?
If it was not zero, lock was set, so loop
I return to caller; critical region entered

I store a 0 in lock
I return to caller

Figure 2-36. Entering and leaving a critical section using the XCHG instruction.
Modern Operating Systems, Tanenbaum & Bos

Easy solution 1 (for a single core OS)

- Provide an enter_region system call that
 - Enters the kernel
 - Disables interrupts
 - Runs the "try part" of enter_region on the right (don't use the loop)
 - If it doesn't succeed, blocks the thread/process and puts it on a waiting queue. Then run scheduler
 - If it succeeds, enable interrupts and return
- Provide a leave_region system call that
 - Enters the kernel
 - Disables the interrupts
 - Pops out one or more processes/threads from the waiting queue
 - Runs the leave_region bit on the right
 - Enables interrupts and returns
- High overhead

```
enter_region:
    TSL REGISTER, LOCK
    CMP REGISTER, #0
    JNE enter_region
    RET

leave_region:
    MOV LOCK, #0
    RET
```

I copy lock to register and set lock to 1
I was lock zero?
If it was not zero, lock was set, so loop
I return to caller; critical region entered

I store a 0 in lock
I return to caller

```
enter_region:
    MOV REGISTER, #1
    XCHG REGISTER, LOCK
    CMP REGISTER, #0
    JNE enter_region
    RET

leave_region:
    MOV LOCK, #0
    RET
```

I put a 1 in the register
I swap the contents of the register and lock variable
I was lock zero?
If it was not zero, lock was set, so loop
I return to caller; critical region entered

I store a 0 in lock
I return to caller

Figure 2-25. Entering and leaving a critical section using the TSL instruction.
Modern Operating Systems, Tanenbaum & Bos

Figure 2-36. Entering and leaving a critical section using the XCHG instruction.
Modern Operating Systems, Tanenbaum & Bos

Producer-consumer problem


The producer-consumer problem has the following parts

- A producer: creates information that it adds to a queue. If the queue is full, it waits until there is room.
- A queue (can be a circular buffer) with a limited number of slots (bounded buffer)
- A consumer: waits until there is something in the queue, pulls the first item out and consumes it (ignore how).

The solutions has to observe the following:

- The producer and consumer are executing independently. You cannot assume anything about their speeds.
- The queue data structure must be preserved. A simple solution is to use a critical region around anything that handles the queue.

More general case: there might be multiple producers and consumers, but we will ignore that for now.



Producer-consumer – solution 1

- Solution that assumes only one producer and one consumer
- Both use sleep if they cannot continue adding or removing items from the queue
- The other end uses wakeup if they add or remove items
- Race condition... where?

```
#define N 100
int count = 0;
/* number of slots in the buffer */
/* number of items in the buffer */

void producer(void)
{
    int item;
    while (TRUE) {
        item = produce_item();
        /* repeat forever */
        /* generate next item */
        if (count == N) sleep();
        /* if buffer is full, go to sleep */
        insert_item(item);
        /* put item in buffer */
        count = count + 1;
        /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);
        /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;
    while (TRUE) {
        if (count == 0) sleep();
        /* if buffer is empty, got to sleep */
        item = remove_item();
        /* take item out of buffer */
        count = count - 1;
        /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer);
        /* was buffer full? */
        consume_item(item);
        /* print item */
    }
}
```

Figure 2-27. Producer-consumer problem with a fatal race condition.
Modern Operating Systems, Tanenbaum & Bos

Producer-consumer – solution 1

- Race condition... where?
- A general tool for spotting potential race conditions
 - These are equivalent for this discussion
 - Read-modify-write
 - Observe-decide-act
 - If the three are not done atomically (anybody can modify state between the steps), then there is a chance of a race condition

```
#define N 100
int count = 0;
/* number of slots in the buffer */
/* number of items in the buffer */

void producer(void)
{
    int item;
    while (TRUE) {
        item = produce_item();
        /* repeat forever */
        /* generate next item */
        if (count == N) sleep();
        /* if buffer is full, go to sleep */
        insert_item(item);
        /* put item in buffer */
        count = count + 1;
        /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);
        /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;
    while (TRUE) {
        if (count == 0) sleep();
        /* if buffer is empty, got to sleep */
        item = remove_item();
        /* take item out of buffer */
        count = count - 1;
        /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer);
        /* was buffer full? */
        consume_item(item);
        /* print item */
    }
}
```

Figure 2-27. Producer-consumer problem with a fatal race condition.
Modern Operating Systems, Tanenbaum & Bos

Producer-consumer – solution 1

- Using the tool 1 (producer vs. consumer):
 - "if (count == N) sleep()" is an **observe** (read count), **decide** (compare count to N), and **act** (sleep).
 - There is no protection of the state, so the consumer may remove an item between reading count and the decide step in the producer.
 - The producer is not woken up even if there is room. May wake up the next time the consumer removes an item.
- Using the tool 2 (consumer vs. producer):
 - Observe from the point of the consumer
 - If consumer preempted just before sleep, it might miss a wakeup from producer (count 0->1). Since this is the only chance, consumer never wakes up.

```

#define N 100 // number of slots in the buffer
int count = 0; // number of items in the buffer

void producer(void)
{
    int item;
    while (TRUE) {
        item = produce_item(); // repeat forever
        if (count == N) sleep(); // generate next item
        insert_item(item); // if buffer is full, go to sleep
        count = count + 1; // put item in buffer
        if (count == N) wakeup(consumer); // increment count of items in buffer
    }
}

void consumer(void)
{
    int item;
    while (TRUE) {
        if (count == 0) sleep(); // repeat forever
        item = remove_item(); // if buffer is empty, go to sleep
        count = count - 1; // take item out of buffer
        consume_item(item); // decrement count of items in buffer
        if (count == N - 1) wakeup(producer); // was buffer full?
        print_item(item); // print item
    }
}
    
```

Figure 2-27. Producer-consumer problem with a fatal race condition. Modern Operating Systems, Tanenbaum & Bos

Semaphores

- Locks typically have two states:
 - 0: lock free / released
 - 1: lock taken / acquired
- A more general concept is a semaphore
 - General idea: use an integer to store the number of wakeups. Can be larger than 1!
 - Two operations (similar to acquire and release). Both are atomic:
 - Down:** check if value is larger than 0. If it was 0, sleep / wait. When it is 1 or larger: count down by one.
 - Up:** add one to the semaphore. If there is a waiting process, release it.
- Can make user level / spinning semaphores
 - Same problem as with locks etc
- More useful if the waiting process is blocked and put on a queue (atomically) with the help of the operating system.

Pthreads semaphores

```

#include <semaphore.h>

int sem_init(sem_t *sem, int pshared, unsigned int value);

int sem_post(sem_t *sem);
int sem_wait(sem_t *sem);
int sem_trywait(sem_t *sem);

int sem_timedwait(sem_t *sem, const struct timespec *abs_timeout);

From manpage of sem_init, sem_post, sem_wait
    
```

Semaphores

- Solving producer-consumer using semaphores
- To understand how it works:
 - The mutex semaphore is used to protect the queue datastructure
 - Note empty=N and full=0
 - Try looking at two cases first
 - 1) Producer running until the queue is full
 - 2) Consumer running until it blocks

```

#define N 100 // number of slots in the buffer
typedef int semaphore; // semaphores are a special kind of int
semaphore mutex = 1; // controls access to critical region
semaphore empty = N; // counts empty buffer slots
semaphore full = 0; // counts full buffer slots


void producer(void)
{
    int item;
    while (TRUE) {
        item = produce_item(); // TRUE is the constant 1
        down(&empty); // generate something to put in buffer
        down(&mutex); // decrement empty count
        insert_item(item); // enter critical region
        up(&full); // put new item in buffer
        up(&mutex); // leave critical region
    }
}

void consumer(void)
{
    int item;
    while (TRUE) {
        down(&full); // infinite loop
        down(&mutex); // decrement full count
        item = remove_item(); // enter critical region
        up(&empty); // take item from buffer
        consume_item(item); // leave critical region
        up(&empty); // increment count of empty slots
    }
}
    
```

Figure 2-28. The producer-consumer problem using semaphores. Modern Operating Systems, Tanenbaum & Bos

Readers and writers problem


- #r >= 1 readers trying to read state
- #w >= 1 writers trying to write state
- Special cases where #r or #w is 1.
- Rules:
 - If no writers are active, then multiple readers can be active at the same time (no changes to the state)
 - If a writer is changing state, then we should not allow any readers (observe inconsistent state) or any other writers (inconsistent updates to state)
- Can get better performance as multiple readers can be serviced at the same time.
- Example from the book: airline reservation database.



Readers and writers problem

Need to keep track of the following invariants

- Number of readers (#r):
 - 0..R if #w == 0, 0 if #w > 0
- Number of writers:
 - 0..1 if #r == 0, 0 if #r > 0



One solution using semaphores

Note the asymmetry

- Writer(s) lock the entire db so only one writer can be inside at any point in time, and it also locks out readers
- The first reader locks down the database on behalf of other readers
 - Releases the mutex semaphore to let other readers in
 - The last reader to exit releases the db
- Note: there is a fairness issue with this solution. A continuous stream of readers may keep the writers locked out indefinitely.

```

typedef int semaphore; // use your imagination
semaphore mutex = 1; // controls access to rc
semaphore db = 1; // controls access to the database
int rc = 0; // # of processes reading or wanting to

void reader(void)
{
    while (TRUE) {
        down(&mutex); // repeat forever
        rc = rc + 1; // get exclusive access to rc
        if (rc == 1) down(&db); // one reader more now
        up(&mutex); // if this is the first reader
        read_data_base(); // release exclusive access to rc
        down(&mutex); // access the data
        rc = rc - 1; // get exclusive access to rc
        if (rc == 0) up(&db); // one reader fewer now
        up(&mutex); // if this is the last reader
        use_data_read(); // release exclusive access to rc
    }
}

void writer(void)
{
    while (TRUE) {
        down(&db); // repeat forever
        write_data_base(); // noncritical region
        up(&db); // get exclusive access
    }
}
    
```

Figure 2-29. A solution to the readers and writers problem. Modern Operating Systems, Tanenbaum & Bos